Saurav Prakash

124, Coordinated Science Laboratory,
University of Illinois Urbana-Champaign, Urbana, IL 61801

✉ sauravp2@illinois.edu
🖥 https://sauravpr.com/

Education

2016 – 2022	**M.S./Ph.D. in Electrical Engineering**.
	University of Southern California, Los Angeles, CA, USA.
	Advisor: Prof. Salman Avestimehr.
2012 – 2016	**B.Tech. in Electrical Engineering**.
	with Minor in Artificial Intelligence in Computer Science and Engineering.
	Indian Institute of Technology Kanpur, Kanpur, UP, India.
	Advisor: Prof. Aditya K. Jagannatham.

Research Interests

- Distributed Optimization and Learning
- Security and Privacy in Data Analytics
- Biological Data Modeling and Analysis
- Information and Coding Theory

Professional Experience

Jul. 2022 – present **Postdoctoral Researcher**.
Coordinated Science Laboratory, University of Illinois Urbana-Champaign, Urbana, IL.
Mentor: Prof. Olgica Milenkovic.

- Machine Unlearning of Federated Clusters
  - Proposed the first framework for the problem of machine unlearning for unsupervised federated clustering (FC)
  - Developed an efficient unlearning mechanism for a customized secure FC framework
  - Designed secure compressed multiset aggregation (SCMA) framework for sparse secure federated learning

- Federated Learning in Hyperbolic Geometry
  - Proposed the first framework of federated learning in Poincaré disc model, a popular model for efficient embedding of large-scale hierarchical data in hyperbolic geometry
  - Developed a novel privacy-preserving pipeline comprising of new quantization and coding theoretic solutions for enabling efficient hyperbolic SVM classifiers in federated setting
  - Manuscript is under preparation

- Tumor Phylogeny Inference
  - Studying the problem of finding a consensus tumor evolution tree from a given set of conflicting input tumor trees
  - Exploring both deterministic and probabilistic approaches for estimating consensus trees

Aug. 2016 – **Graduate Research Assistant**.

May 2022 *vITAL Lab, University of Southern California*, Los Angeles, CA.

Mentor: Prof. Salman Avestimehr.

○ Resource-Constrained Federated Learning of Large Models
  - Proposed a novel problem formulation of federated learning where no clients are capable of training a full model
  - Proposed PriSM, a probabilistic sampling for sub-models at clients to preserve full-model capacity and achieve memory-efficiency

○ Lottery Aware Sparse Federated Learning
  - Established the reasons behind accuracy drop in training of sparse models in FL settings
  - Proposed FLASH, a unified sparse learning framework to make the server win a lottery in terms of yielding a sparse sub-model
  - Developed specialized extensions for model personalization and heterogenous models

○ Secure and Fault Tolerant Decentralized Learning
  - Considered general privacy-preserving but faulty federated learning setting with non-IID data across clients
  - Proposed DiverseFL, a novel TEE assisted sampling based approach that applies per client criteria for mitigating faults in the general federated learning setting
  - Via extensive experiments, demonstrated that compared to prior approaches, DiverseFL performs much better, almost achieving the optimal model performance

○ Byzantine-Robust Decentralized (Serverless) Learning
  - Considered the problem of Byzantine mitigation in the decentralized learning setting without any central coordinator
  - Developed Basil, a fast and computationally efficient Byzantine-robust algorithm leveraging a sequential, memory assisted and performance criteria for training over a ring
  - Demonstrated numerically that under different Byzantine fault settings, Basil provides up to ∼16% higher test accuracy when compared with prior methods

○ Coded Computing for Hierarchical Distributed Learning at the Edge
  - Formulated a hierarchical gradient aggregation problem for machine learning from data available at the client nodes by leveraging reliable helper nodes for collecting updates
  - Proposed two unique coded computing strategies – aligned repetition coding (ARC), aligned MDS coding (AMC) – for mitigating straggling links from clients to helpers

○ Coded Computing for Large-scale Distributed Learning
  - Formulated a tree gradient coding framework and proposed CodedReduce scheme for fast and robust gradient aggregation in distributed learning
  - CodedReduce combines advantages of communication efficiency of Ring-AllReduce and straggler resiliency of Gradient Coding for minimizing the overall training latency
  - In experiments over Amazon EC2, CodedReduce provides gains of up to $31\times$ in the overall execution time over prior approaches for distributed learning

- Coded Computing for Large-scale Graph Processing
  - Proposed a distributed computing framework for graph analytics based on MapReduce
  - Characterized the optimal trade-off between Map computations and Shuffle load for the Erdos-Renyi model
  - Developed and implemented a coded distributed implementation of the PageRank algorithm using Amazon EC2, demonstrating gains of up to 50% over the naive PageRank
  - Developed coding schemes for three other popular random graph models – random bi-partite model, stochastic block model, and power law model
- Coded Computing for Large-scale Matrix Multiplication in Heterogeneous Settings
  - Proposed a two-step alternative formulation to the problem of minimizing the expected run-time in distributed matrix-vector multiplication in heterogeneous clusters
  - Developed a scalable method – Heterogeneous Coded Matrix Multiplication (HCMM) – for reliable matrix multiplication on cloud clusters with stragglers
  - Proved the asymptotic optimality of HCMM
  - Implemented HCMM using Amazon EC2, demonstrating gains of up to 61% over benchmark schemes

May – Aug. **Graduate Technical Intern**.
2018 & 2019 *Intel Labs*, Santa Clara, CA.

Mentors: Sagar Dhakal, Nageen Himayat, Shilpa Talwar.

- Coded Computing for Federated Learning in Multi-access Edge Computing (MEC) networks
  - Proposed CodedFedL for injecting coding redundancy into federated learning with non-IID data for mitigating stragglers and minimizing training time in MEC networks
  - Developed a tractable approach for minimizing deadline time
  - Analyzed the convergence rate and iteration complexity of CodedFedL
  - Demonstrated gains of up to $15\times$ in comparison to benchmark schemes in practice

Jun. 2021 – **Applied Scientist Intern**.
Aug. 2021 *Alexa AI, Amazon*, Cambridge, MA.

Mentors: Clement Chung, Christophe Dupuy, Rahul Gupta, Leo Long, Tanya Roosta.

- Federated learning with Heterogeneous Model Architectures
  - Developed various strategies for efficient federated learning from edge users
  - Explored novel methods to enable federated learning with heterogeneous model architectures at the edge users

May 2015 – **International Visiting Student**.
Jul. 2015 *IUSSTF-Viterbi Program*, Los Angeles, CA.

Mentor: Prof. Salman Avestimehr.

- Towards Faster Algorithms for Processing Large Data on Graphs
  - Studied spectral graph theory and its application in signal processing of graph data – cut-off frequency, optimal sampling and bandlimited interpolation
  - Explored existing semi-supervised and active learning methods for data on graphs
  - Proposed Random Jump model based on Graph Laplacian for low complexity sampling

2013 – 2016 **Undergraduate Research Assistant**.
MWN Group, IIT Kanpur, Kanpur, India.

Mentor: Prof. Aditya K. Jagannatham.

○ Scheduling for Efficient Utilization of Time Resource in Wireless Networks
  - Worked on the problem of user scheduling for efficient wireless resource utilization, under resource allocation fairness constraints
  - Proposed two opportunistic schemes for scheduling users in a time slotted system with wireless Rayleigh-fading channel
  - Simulations predicted stochastically improved performance compared to Round Robin scheme alongside satisfaction of any arbitrary time/resource allocation fairness constraints

May 2013 – **Undergraduate Research Intern**.
Jul. 2013 *Summer Undergraduate Research Grant for Excellence (SURGE), IIT Kanpur*, Kanpur, India.

Mentor: Prof. Aditya K. Jagannatham.

○ Channel Estimation and Capacity in MIMO Wireless Communication Systems
  - Studied capacity lower bound for a MIMO system obeying Block-Fading law using LMMSE estimator for channel estimation at the receiver
  - Simulated existing algorithms designed for obtaining near optimal number of transmit antennas for optimizing the capacity lower bound

Selected Honors and Awards

2023 Institute for Genomic Biology (IGB) Fellowship.

2021 Qualcomm Innovation Fellowship.

2019 Qualcomm Innovation Fellowship Finalist.

2019 Most Novel Research Project Award, EE-599 (Deep Learning course, USC).

2016 USC Annenberg PhD Fellowship.

2016 Princeton Gordon Wu PhD Fellowship Offer.

2015 Viterbi-India Internship.

2014 Summer Undergraduate Research Grant for Excellence at IIT Kanpur.

2015 Shri Singhasan Singh Scholarship at IIT Kanpur.

2015 Institution of Engineering and Technology (IET) Scholarship.

Publications

Link to Google Scholar.

(* denotes joint authorship).

Conference/Workshop Proceedings

C8 C. Pan*, J. Sima*, **S. Prakash***, V. Rana, O. Milenkovic, "Machine Unlearning of Federated Clusters ," *The Eleventh International Conference on Learning Representations (ICLR)*, Jan. 2023.

C7 **S. Prakash***, A. Reisizadeh*, R. Pedarsani, S. Avestimehr, "Hierarchical coded gradient aggregation for learning at the edge," in *Proceedings of IEEE International Symposium on Information Theory (ISIT)*, Aug. 2020 .

C6 S. Dhakal, **S. Prakash**, Y. Yona, S. Talwar, N. Himayat, "Coded federated learning," in *Proceedings of IEEE Globecom Workshops (GC Wkshps)*, Mar. 2020.

C5 S. Kundu*, **S. Prakash***, H. Akrami, P. Beerel, K. Chugg, "pSConv: A pre-defined sparse kernel based convolution for deep CNNs," in *Proceedings of IEEE 57th Annual Allerton Conference on Communication, Control, and Computing (Allerton)*, Dec. 2019.

C4 S. Dhakal*, **S. Prakash***, Y. Yona, S. Talwar, N. Himayat, "Coded computing for distributed machine learning in wireless edge network," in *Proceedings of IEEE 90th Vehicular Technology Conference (VTC2019-Fall)*, Nov. 2019.

C3 A. Reisizadeh*, **S. Prakash***, R. Pedarsani, S. Avestimehr, "Tree gradient coding," in *Proceedings of IEEE International Symposium on Information Theory (ISIT)*, Sep. 2019.

C2 **S. Prakash***, A. Reisizadeh*, R. Pedarsani, S. Avestimehr, "Coded computing for distributed graph analytics," in *Proceedings of IEEE International Symposium on Information Theory (ISIT) Conference*, Aug. 2018.

C1 A. Reisizadeh, **S. Prakash**, R. Pedarsani, S. Avestimehr, "Coded computation over heterogeneous clusters," in *Proceedings of IEEE International Symposium on Information Theory (ISIT) Conference*, Aug. 2017.

Journal Papers

J5 A. R. Elkordy, **S. Prakash**, S. Avestimehr, "Basil: A fast and Byzantine-resilient approach for decentralized training," *IEEE Journal on Selected Areas in Communications*, volume 40, issue 9, pages 2694–2716, Sep. 2022. *Part of it was presented at the NeurIPS Workshop on Privacy in Machine Learning, 2021*.

J4 **S. Prakash**, S. Dhakal, M. Akdeniz, Y. Yona, S. Talwar, S. Avestimehr, N. Himayat, "Coded computing for low-latency federated learning over wireless edge networks," *IEEE Journal on Selected Areas in Communications*, volume 39, issue 1, pages 233–250, Jan. 2021**.** *Was partly presented at the FL-ICML Workshop on User Privacy and Data Confidentiality, 2020*.

J3 A. Reisizadeh*, **S. Prakash***, R. Pedarsani, S. Avestimehr, "CodedReduce: A fast and robust framework for gradient aggregation in distributed learning," *IEEE/ACM Transactions on Networking*, volume 30, issue 1, pages 148–161, Feb. 2022**.** *Was partly presented at the ICML Workshop on Coding Theory for Large-Scale Machine Learning (CodML), 2019*.

J2 **S. Prakash***, A. Reisizadeh*, R. Pedarsani, S. Avestimehr, "Coded computing for distributed graph analytics," *IEEE Transactions on Information Theory*, volume 66, issue 10, pages 6534–6554, Oct. 2020 .

J1 A. Reisizadeh, **S. Prakash**, R. Pedarsani, S. Avestimehr, "Coded computation over heterogeneous clusters," *IEEE Transactions on Information Theory*, volume 65, issue 7, pages 4227–4242, Jul. 2019 .

Preprints

P3 **S. Prakash**, H. Hashemi, Y. Wang, M. Annavaram, S. Avestimehr, "Secure and Fault Tolerant Decentralized Learning," *arXiv:2109.07706*, Sep. 2022. *Partly presented at the Enclaved AI/ML Workshop 2021, Private AI Research Institute*.

P2 Y. Niu*, **S. Prakash***, S. Kundu, S. Lee, S. Avestimehr, "Federated Learning of Large Models at the Edge via Principal Sub-Model Training," *arXiv:2208.13141*, Oct. 2022. *Partly presented at the NeurIPS Workshop on Recent Advances and New Challenges in Federated Learning (FL-NeurIPS), 2022*.

P1 S. Babakniya, S. Kundu, **S. Prakash**, Y. Niu, S. Avestimehr, "Federated Sparse Training: Lottery Aware Model Compression for Resource Constrained Edge," *arXiv:2208.13092*, Oct. 2022. *Partly presented at the NeurIPS Workshop on Recent Advances and New Challenges in Federated Learning (FL-NeurIPS), 2022*.

Patents

2022 **S. Prakash**, S. Dhakal, Y. Yona, N. Himayat, S. Talwar, "Technologies for distributing iterative computations in heterogeneous computing environments," *US Patent 11,423,254*.

2022 **S. Prakash**, S. Dhakal, Y. Yona, N. Himayat, S. Talwar, "Technologies for distributing gradient descent computation in a heterogeneous multi-access edge computing (MEC) networks," *US Patent 11,244,242*.

2021 M. R. Akdeniz, A. Anand, N. Himayat, A. S. Avestimehr, R. Balakrishnan, P. Bhardwaj, J. Choi, Y.-S. Choi, S. Dhakal, B. G. Edwards, **S. Prakash**, A. Solomon, S. Talwar, Y. E. Yona, "Systems and methods for distributed learning for wireless edge dynamics," *App. No. PCT/US2020/067068*.

Professional Service

2017–present **Invited Journal Reviewer**.
- IEEE Journal on Selected Areas in Communications (16)
- IEEE Transactions on Communications (12)
- IEEE Journal on Selected Areas in Information Theory (5)
- IEEE Transactions on Information Theory (4)
- IEEE/ACM Transactions on Networking (2)
- IEEE Transactions on Information Forensics and Security (2)
- IEEE Transactions on Mobile Computing (2)
- IEEE Internet of Things Magazine (1)
- IEEE Open Journal of the Communications Society (3)
- IEEE Communications Letters (1)

2017–present **Invited Conference/Workshop Reviewer**.
- Knowledge Discovery and Data Mining (KDD) (3)
- Artificial Intelligence and Statistics (AISTATS) (7)
- IEEE International Symposium on Information Theory (ISIT) (8)
- IEEE Information Theory Workshop (ITW) (2)
- Resource-Constrained Learning in Wireless Networks (RCLWN) @ MLSys 2023 (1)
- Federated Learning for Natural Language Processing (FL4NLP) @ ACL 2022 (2)

Mentorship Experience

2021 **Mentor**, *Graduate Application Mentorship Program (GradAMP)*, USC.

2013 – 2016 **Student Guide**, *Counselling Service*, IIT Kanpur.

2013 – 2015 **Ambassador Caller**, *Alumni Contact Program*, IIT Kanpur.

2013 – 2014 **Student Secretary**, *Fine Arts Club*, IIT Kanpur.

Selected Coursework

Algorithms and Artificial Intelligence.

- Data Structure and Algorithms
- Deep Learning
- Fundamentals of Computing
- Machine Learning for Computer Vision

Communication and Signal Processing.

- Information Theory and Compression
- Error Correcting Codes
- Topics in Cryptography and Coding
- Digital Communication and Coding Systems

Optimization and Statistics.

- Optimization: Theory and Algorithms
- Introduction to Mathematical Statistics
- High Dimensional Statistics and Big Data Problems
- Network Flows and Combinatorial Optimization

Mathematics.

- Theory of Probability
- Real Analysis
- Fundamentals of Modern Algebra
- Partial Differential Equations