

Saurav Prakash

Ming Hsieh Department of Electrical and Computer Engineering,
University of Southern California, Los Angeles, CA 90089

✉ sauravpr@usc.edu
🌐 <https://sauravpr.com/>

Education

2016 – **Ph.D. in Electrical Engineering.**

Present *University of Southern California*, Los Angeles, CA, USA.

Advisor: Prof. Salman Avestimehr.

2012 – 2016 **B.Tech. in Electrical Engineering.**

with Minor in Artificial Intelligence in Computer Science and Engineering.

Indian Institute of Technology Kanpur, Kanpur, UP, India.

Advisor: Prof. Aditya K. Jagannatham.

Research Interests

- Security and Privacy in Machine Learning
- Efficient and Robust Federated Learning
- Large-Scale Decentralized Serverless Training
- Coded Distributed Computing
- Information and Coding Theory

Research Overview

I have worked extensively towards holistically addressing real-world bottlenecks in large-scale distributed computing, including both *large-scale cloud computing* and *decentralized machine learning*. My contributions involve novel solutions that address multiple challenges in both these broad domains, bringing new concepts from Shannon's information theory, coding theory, communication topology design, data privacy, and optimization theory.

Professional Experience

Aug. 2016 – **Graduate Research Assistant.**

Present *vITAL Lab, University of Southern California*, Los Angeles, CA.

Mentor: Prof. Salman Avestimehr.

- Byzantine Robust Federated Learning
 - Considered general Byzantine federated learning setting with non-IID data across clients
 - Proposed DiverseFL, a novel sampling based approach that applies per client criteria for mitigating Byzantines in the general federated learning setting
 - Demonstrated via extensive experiments that compared to prior approaches, DiverseFL performs much better, almost achieving the optimal model performance

- Byzantine Robust Decentralized (Serverless) Learning
 - Considered the problem of Byzantine mitigation in the decentralized learning setting without any central coordinator
 - Developed Basil, a fast and computationally efficient Byzantine robust algorithm leveraging a sequential, memory assisted and performance criteria for training over a ring
 - Demonstrated numerically that Basil provides up to $\sim 16\%$ higher test accuracy when compared with prior methods, under different Byzantine fault settings
- Coded Computing for Hierarchical Distributed Learning at the Edge
 - Formulated a hierarchical gradient aggregation problem for machine learning from data available at the client nodes by leveraging reliable helper nodes for collecting updates
 - Proposed two unique coded computing strategies – aligned repetition coding (ARC), aligned MDS coding (AMC) – for mitigating straggling links from clients to helpers
- Coded Computing for Large-scale Distributed Learning
 - Formulated a tree gradient coding framework and proposed CodedReduce scheme for fast and robust gradient aggregation in distributed learning
 - CodedReduce combines advantages of communication efficiency of Ring-AllReduce and straggler resiliency of Gradient Coding for minimizing the overall training latency
 - In experiments over Amazon EC2, CodedReduce achieves gains of up to $31\times$ in the overall execution time over prior approaches for distributed learning
- Coded Computing for Large-scale Graph Processing
 - Proposed a distributed computing framework for graph analytics based on MapReduce
 - Characterized the optimal trade-off between Map computations and Shuffle load for the Erdos-Renyi model
 - Developed and implemented a coded distributed implementation of the PageRank algorithm using Amazon EC2, demonstrating gains of up to 50% over the naive PageRank
 - Developed coding schemes for three other popular random graph models – random bi-partite model, stochastic block model, and power law model
- Coded Computing for Large-scale Matrix Multiplication in Heterogeneous Settings
 - Proposed a two-step alternative formulation to the problem of minimizing the expected run-time in distributed matrix-vector multiplication in heterogeneous clusters
 - Developed a scalable method – Heterogeneous Coded Matrix Multiplication (HCMM) – for reliable matrix multiplication on cloud clusters with stragglers
 - Proved the asymptotic optimality of HCMM
 - Implemented HCMM using Amazon EC2, demonstrating gains of up to 61% over benchmark schemes

May – Aug. **Graduate Technical Intern.**

2018 & 2019 *Intel Labs*, Santa Clara, CA.

Mentors: Sagar Dhakal, Nageen Himayat, Shilpa Talwar.

- Coded Computing for Federated Learning in Multi-access Edge Computing (MEC) networks
 - Proposed CodedFedL for injecting coding redundancy into federated learning with non-IID data for mitigating stragglers and minimizing up training time in MEC networks
 - Developed a tractable approach for minimizing deadline time
 - Analyzed the convergence rate and iteration complexity of CodedFedL
 - Demonstrated gains of up to $15\times$ in comparison to benchmark schemes in practice

Jun. 2021 – **Applied Scientist Intern.**

Aug. 2021 *Alexa AI, Amazon*, Cambridge, MA.

Mentors: Clement Chung, Christophe Dupuy, Rahul Gupta, Leo Long, Tanya Roosta.

- Federated learning with Heterogeneous Model Architectures
 - Developed various strategies for efficient federated learning from edge users
 - Explored novel methods to enable federated learning with heterogeneous model architectures at the edge users

May 2015 – **International Visiting Student.**

Jul. 2015 *IUSSTF-Viterbi Program*, Los Angeles, CA.

Mentor: Prof. Salman Avestimehr.

- Towards Faster Algorithms for Processing Large Data on Graphs
 - Studied spectral graph theory and its application in signal processing of graph data – cut-off frequency, optimal sampling and bandlimited interpolation
 - Explored existing semi-supervised and active learning methods for data on graphs
 - Proposed a Random Jump model based on Graph Laplacian for sampling with low time complexity

2013 – 2016 **Undergraduate Research Assistant.**

MWN Group, IIT Kanpur, Kanpur, India.

Mentor: Prof. Aditya K. Jagannatham.

- Scheduling for Efficient Utilization of Time Resource in Wireless Networks
 - Worked on the problem of user scheduling for efficient wireless resource utilization, under resource allocation fairness constraints
 - Proposed two opportunistic schemes for scheduling users in a time slotted system with wireless Rayleigh-fading channel
 - Simulations predicted stochastically improved performance compared to Round Robin scheme alongside satisfaction of any arbitrary time resource allocation fairness constraints

May 2013 – **Undergraduate Research Intern.**

Jul. 2013 *Summer Undergraduate Research Grant for Excellence (SURGE)*, IIT Kanpur, Kanpur, India.

Mentor: Prof. Aditya K. Jagannatham.

- Channel Estimation and Capacity in MIMO Wireless Communication Systems
 - Studied capacity lower bound for a MIMO system obeying Block-Fading law using LMMSE estimator for channel estimation at the receiver
 - Simulated existing algorithms designed for obtaining near optimal number of transmit antennas for optimizing the capacity lower bound

Selected Honors and Awards

2021 Qualcomm Innovation Fellowship.

2019 Qualcomm Innovation Fellowship Finalist.

2019 Most Novel Research Project Award, EE-599 (Deep Learning course, USC).

2016 USC Annenberg PhD Fellowship.

- 2016 Princeton Gordon Wu PhD Fellowship Offer.
- 2015 Viterbi-India Internship.
- 2014 Summer Undergraduate Research Grant for Excellence at IIT Kanpur.
- 2015 Shri Singhasan Singh Scholarship at IIT Kanpur.
- 2015 Institution of Engineering and Technology (IET) Scholarship.

Publications

Link to Google Scholar.

(* denotes joint authorship).

Preprints

- P2 **S. Prakash**, H. Hashemi, Y. Wang, M. Annavaram, S. Avestimehr, “Byzantine resilient federated learning with heterogeneous data distribution,” Jul. 2021. *Partly presented at the Enclaved AI/ML Workshop 2021, Private AI Research Institute.*
- P1 A. R. Elkordy, **S. Prakash**, S. Avestimehr, “Basil: A fast and Byzantine-resilient approach for decentralized training,” Sep. 2021. *Part of it to be presented at the NeurIPS Workshop on Privacy in Machine Learning, 2021.*

Journal Papers

- J4 **S. Prakash**, S. Dhakal, M. Akdeniz, Y. Yona, S. Talwar, S. Avestimehr, N. Himayat, “Coded computing for low-latency federated learning over wireless edge networks,” *IEEE Journal on Selected Areas in Communications*, volume 39, issue 1, pages 233–250, Jan. 2021. *Was partly presented at the FL-ICML Workshop on User Privacy and Data Confidentiality, 2020.*
- J3 A. Reisizadeh*, **S. Prakash***, R. Pedarsani, S. Avestimehr, “Coded Reduce: A fast and robust framework for gradient aggregation in distributed learning,” to appear in the *IEEE/ACM Transactions on Networking*.
- J2 **S. Prakash***, A. Reisizadeh*, R. Pedarsani, S. Avestimehr, “Coded computing for distributed graph analytics,” *IEEE Transactions on Information Theory*, volume 66, issue 10, pages 6534–6554, Oct. 2020 .
- J1 A. Reisizadeh, **S. Prakash**, R. Pedarsani, S. Avestimehr, “Coded computation over heterogeneous clusters,” *IEEE Transactions on Information Theory*, volume 65, issue 7, pages 4227–4242, Jul. 2019 .

Conference/Workshop Proceedings

- C7 **S. Prakash***, A. Reisizadeh*, R. Pedarsani, S. Avestimehr, “Hierarchical coded gradient aggregation for learning at the edge,” in *Proceedings of IEEE International Symposium on Information Theory (ISIT)*, Aug. 2020 .
- C6 S. Dhakal, **S. Prakash**, Y. Yona, S. Talwar, N. Himayat, “Coded federated learning,” in *Proceedings of IEEE Globecom Workshops (GC Wkshps)*, Mar. 2020.
- C5 S. Kundu*, **S. Prakash***, H. Akrami, P. Beerel, K. Chugg, “pSConv: A pre-defined sparse kernel based convolution for deep CNNs,” in *Proceedings of IEEE 57th Annual Allerton Conference on Communication, Control, and Computing (Allerton)*, Dec. 2019.

- C4 S. Dhakal*, **S. Prakash***, Y. Yona, S. Talwar, N. Himayat, “Coded computing for distributed machine learning in wireless edge network,” in *Proceedings of IEEE 90th Vehicular Technology Conference (VTC2019-Fall)*, Nov. 2019.
- C3 A. Reisizadeh*, **S. Prakash***, R. Pedarsani, S. Avestimehr, “Tree gradient coding,” in *Proceedings of IEEE International Symposium on Information Theory (ISIT)*, Sep. 2019.
- C2 **S. Prakash***, A. Reisizadeh*, R. Pedarsani, S. Avestimehr, “Coded computing for distributed graph analytics,” in *Proceedings of IEEE International Symposium on Information Theory (ISIT) Conference*, Aug. 2018.
- C1 A. Reisizadeh, **S. Prakash**, R. Pedarsani, S. Avestimehr, “Coded computation over heterogeneous clusters,” in *Proceedings of IEEE International Symposium on Information Theory (ISIT) Conference*, Aug. 2017.

Selected Talks

- Jul. 2021 **TEE-GPU Cooperative Learning: Privacy and Security Without the Price**, *Presentation, Enclaved AI/ML Workshop 2021*, Private AI Research Institute.
- May 2021 **Federated deep learning: On-device learning for CV and NLP**, *Finalist Team Presentation, Qualcomm Innovation Fellowship 2021*, Qualcomm.
- Apr 2021 **Trustworthy and Scalable Federated Learning**, *CCF Advanced Disciplines Lecture*, Institute of Computing Technology, Chinese Academy of Sciences.
- Jul. 2020 **Coded Computing for Federated Learning at the Edge**, *Presentation, FL Workshop on User Privacy and Data Confidentiality*, International Conference on Machine Learning.

Patents

- 2021 M. R. Akdeniz, A. Anand, N. Himayat, A. S. Avestimehr, R. Balakrishnan, P. Bhardwaj, J. Choi, Y.-S. Choi, S. Dhakal, B. G. Edwards, **S. Prakash**, A. Solomon, S. Talwar, Y. E. Yona, “Systems and methods for distributed learning for wireless edge dynamics,” *App. No. PCT/US2020/067068*.
- 2019 **S. Prakash**, S. Dhakal, Y. Yona, N. Himayat, S. Talwar, “Technologies for distributing iterative computations in heterogeneous computing environments,” *US Patent App. 16/368,716*.
- 2019 **S. Prakash**, S. Dhakal, Y. Yona, N. Himayat, S. Talwar, “Technologies for distributing gradient descent computation in a heterogeneous multi-access edge computing (MEC) networks,” *US Patent App. 16/235,682*.

Professional Service

- 2017–2021 **Invited Journal Reviewer.**
 - IEEE Journal on Selected Areas in Communications
 - IEEE Transactions on Information Theory
 - IEEE Journal on Selected Areas in Information Theory
 - IEEE Transactions on Communications
- 2017–2021 **Invited Conference/Workshop Reviewer.**
 - IEEE International Symposium on Information Theory (ISIT)
 - IEEE Information Theory Workshop (ITW)

Mentorship Experience

- 2021 **Mentor**, *Graduate Application Mentorship Program (GradAMP)*, USC.
- 2013 – 2016 **Student Guide**, *Counselling Service*, IIT Kanpur.
- 2013 – 2015 **Ambassador Caller**, *Alumni Contact Program*, IIT Kanpur.
- 2013 – 2014 **Student Secretary**, *Fine Arts Club*, IIT Kanpur.

Selected Coursework

Algorithms and Artificial Intelligence.

- Data Structure and Algorithms
- Deep Learning
- Fundamentals of Computing
- Machine Learning for Computer Vision

Communication and Signal Processing..

- Information Theory and Compression
- Error Correcting Codes
- Topics in Cryptography and Coding
- Digital Communication and Coding Systems

Optimization and Statistics.

- Optimization: Theory and Algorithms
- High Dimensional Statistics and Big Data Problems
- Network Flows and Combinatorial Optimization
- Introduction to Mathematical Statistics

Mathematics.

- Partial Differential Equations
- Real Analysis
- Theory of Probability
- Fundamentals of Modern Algebra