

Saurav Prakash

CONTACT INFORMATION

EEB 514, 3740 McClintock Ave,
Electrical and Computer Engineering,
University of Southern California,
Los Angeles, CA 90089.
Email: sauravpr@usc.edu
Phone: +1 2132459100

<https://sauravpr.com/>
[Google Scholar](#)
[LinkedIn](#)

RESEARCH INTERESTS

Secure and Privacy-Preserving Machine Learning, Federated Learning, Large-Scale Serverless Training, Coded Distributed Computing, Information and Coding Theory

EDUCATION

Doctor of Philosophy in Electrical and Computer Engineering
University of Southern California
Advisor – [Prof. Salman Avestimehr](#)

Aug 2016 – present

Bachelor of Technology in Electrical Engineering
Minor in Artificial Intelligence in Computer Science and Engineering
Indian Institute of Technology Kanpur
Advisor – [Prof. Aditya K. Jagannatham](#)

Jul 2012 – May 2016

PUBLICATIONS

Preprints and Journal Articles

- Saurav Prakash, Hanieh Hashemi, Yongqin Wang, Murali Annavaram, Salman Avestimehr. “[Byzantine resilient federated learning with heterogeneous data distribution](#)”, arXiv preprint arXiv:2010.07541, 2021.
- Ahmed Roushdy Elkordy, Saurav Prakash, Salman Avestimehr. “[Basil: A fast and Byzantine-resilient approach for decentralized training](#),” arXiv preprint arXiv:2109.07706, 2021.
- Saurav Prakash, Sagar Dhakal, Mustafa Akdeniz, Yair Yona, Shilpa Talwar, Salman Avestimehr Nageen Himayat. “[Coded computing for low-latency federated learning over wireless edge networks](#),” IEEE Journal on Selected Areas in Communications, 2021.
- Saurav Prakash, Sagar Dhakal, Mustafa Akdeniz, Salman Avestimehr, Nageen Himayat. “[Coded computing for federated learning at the edge](#),” arXiv preprint arXiv:2007.03273, 2020. Presented at the FL-ICML Workshop on User Privacy and Data Confidentiality, 2020.
- Amirhossein Reisizadeh*, Saurav Prakash*, Ramtin Pedarsani, Salman Avestimehr. “[Coded Reduce: A fast and robust framework for gradient aggregation in distributed learning](#),” IEEE/ACM Transactions on Networking, 2021. Partly presented at the ICML Workshop on Coding Theory for Large-Scale Machine Learning, 2019.
- Saurav Prakash*, Amirhossein Reisizadeh*, Ramtin Pedarsani, Salman Avestimehr. “[Coded computing for distributed graph analytics](#),” IEEE Transactions on Information Theory, 2020.
- Amirhossein Reisizadeh, Saurav Prakash, Ramtin Pedarsani, Salman Avestimehr. “[Coded computation over heterogeneous clusters](#),” IEEE Transactions on Information Theory, 2019.

Conference and Workshop Proceedings

- Saurav Prakash*, Amirhossein Reisizadeh*, Ramtin Pedarsani, Salman Avestimehr. “[Hierarchical coded gradient aggregation for learning at the edge](#),” IEEE International Symposium on Information Theory, 2020.
- Sagar Dhakal, Saurav Prakash, Yair Yona, Shilpa Talwar Nageen Himayat. “[Coded federated learning](#),” IEEE Globecom Workshop on Wireless Edge Intelligence, 2019.
- Souvik Kundu*, Saurav Prakash*, Haleh Akrami, Peter Beerel, Keith Chugg. “[A pre-defined sparse kernel based convolution for deep CNNs](#),” IEEE Allerton Conference on Communication, Control, and Computing, 2019.
- Sagar Dhakal*, Saurav Prakash*, Yair Yona, Shilpa Talwar Nageen Himayat. “[Coded computing for distributed machine learning in wireless edge network](#),” IEEE VTC Fall Workshop on Vehicular Information Services for the Internet of Things, 2019.
- Amirhossein Reisizadeh*, Saurav Prakash*, Ramtin Pedarsani, Salman Avestimehr. “[Tree gradient coding](#),” IEEE International Symposium on Information Theory, 2019.
- Saurav Prakash*, Amirhossein Reisizadeh*, Ramtin Pedarsani, Salman Avestimehr. “[Coded computing for distributed graph analytics](#),” IEEE International Symposium on Information Theory, 2018.

* denotes equal contribution

	<ul style="list-style-type: none"> Amirhossein Reisizadeh, Saurav Prakash, Ramtin Pedarsani, Salman Avestimehr. “Coded computation over heterogeneous clusters,” IEEE International Symposium on Information Theory, 2017.
PATENTS	<ul style="list-style-type: none"> Saurav Prakash, Sagar Dhakal, Yair Yona, Nageen Himayat and Shilpa Talwar. “Technologies for distributing iterative computations in heterogeneous computing environments,” US Patent App. 16/368,716. Saurav Prakash, Sagar Dhakal, Yair Yona, Nageen Himayat and Shilpa Talwar. “Technologies for distributing gradient descent computation in a heterogeneous multi-access edge computing (MEC) networks,” US Patent App. 16/235,682.
AWARDS AND ACHIEVEMENTS	<ul style="list-style-type: none"> Qualcomm Innovation Fellowship 2021 Qualcomm Innovation Fellowship Finalist 2019 Most Novel Research Project Award, EE-599 (Deep Learning course, USC) 2019 USC Annenberg PhD Fellowship 2016-2020 Princeton Gordon Wu PhD Fellowship (Declined) 2016 Viterbi-India Internship 2015 Summer Undergraduate Research Grant for Excellence at IIT Kanpur 2014 Shri Singhasan Singh Scholarship at IIT Kanpur 2015 Institution of Engineering and Technology (IET) Scholarship 2015
PROFESSIONAL EXPERIENCE	<p>Graduate Research Assistant <i>Sep 2016 – present</i></p> <p><i>vITAL Lab, University of Southern California, Los Angeles, CA.</i></p> <p>Mentor – Prof. Salman Avestimehr</p> <ul style="list-style-type: none"> Byzantine Robust Federated Learning <ul style="list-style-type: none"> Considered general Byzantine federated learning setting with non-IID data across clients Proposed DiverseFL, a novel sampling based approach that applies per client criteria for mitigating Byzantines in the general federated learning setting Demonstrated via extensive experiments that compared to prior approaches, DiverseFL performs much better, almost achieving the optimal model performance Byzantine Robust Decentralized (Serverless) Learning <ul style="list-style-type: none"> Considered the problem of Byzantine mitigation in the decentralized learning setting without any central coordinator Developed Basil, a fast and computationally efficient Byzantine robust algorithm leveraging a sequential, memory assisted and performance criteria for training over a logical ring Demonstrated numerically that Basil provides up to $\sim 16\%$ higher test accuracy when compared with prior methods, under different Byzantine fault settings Coded Computing for Hierarchical Distributed Learning at the Edge <ul style="list-style-type: none"> Formulated a hierarchical gradient aggregation problem for machine learning from data available at the client nodes by leveraging reliable helper nodes for collecting updates Proposed two unique coded computing strategies – aligned repetition coding (ARC), aligned MDS coding (AMC) – for mitigating straggling links from clients to helpers Coded Computing for Large-scale Distributed Learning <ul style="list-style-type: none"> Formulated a tree gradient coding framework and proposed CodedReduce scheme for fast and robust gradient aggregation in distributed learning CodedReduce combines advantages of communication efficiency of Ring-AllReduce and straggler resiliency of Gradient Coding for minimizing the overall training latency In experiments over Amazon EC2, CodedReduce achieves gains of up to $31\times$ in the overall execution time over prior approaches for distributed learning Coded Computing for Large-scale Graph Processing <ul style="list-style-type: none"> Proposed a distributed computing framework for graph analytics based on MapReduce Characterized the optimal trade-off between Map computations and Shuffle load for the Erdos-Renyi model Developed and implemented a coded distributed implementation of the PageRank algorithm using Amazon EC2, demonstrating gains of up to 50% over the naive PageRank

- Developed coding schemes for three other popular random graph models – random bipartite model, stochastic block model, and power law model
- Coded Computing for Large-scale Matrix Multiplication in Heterogeneous Settings
 - Proposed a two-step alternative formulation to the problem of minimizing the expected run-time in distributed matrix-vector multiplication in heterogeneous clusters
 - Developed a scalable method – Heterogeneous Coded Matrix Multiplication (HCMM) – for reliable matrix multiplication on cloud clusters with stragglers
 - Proved the asymptotic optimality of HCMM
 - Implemented HCMM using Amazon EC2, demonstrating gains of up to 61% over benchmark schemes

Applied Scientist Intern

Summer 2021

Alexa AI, Cambridge, MA.

Mentors – [Clement Chung](#), [Christophe Dupuy](#), [Rahul Gupta](#), [Leo Long](#), [Tanya Roosta](#)

- Federated learning with Heterogeneous Model Architectures
 - Developed various strategies for efficient federated learning from edge users
 - Explored novel methods to enable federated learning with heterogeneous model architectures at the edge users

Graduate Technical Intern

Summers 2018, 2019

Intel Labs, Santa Clara, CA.

Mentors – [Sagar Dhakal](#), [Nageen Himayat](#), [Shilpa Talwar](#)

- Coded Computing for Federated Learning in Multi-access Edge Computing (MEC) networks
 - Proposed the first coded computing framework, CodedFedL, that injects structured coding redundancy into non-linear federated learning for mitigating stragglers and speeding up the training procedure in heterogeneous multi-access edge computing networks
 - Developed a tractable approach for finding optimal coding redundancy and the number of local data points that a client processes during training for minimizing the deadline time
 - Characterized the leakage in data privacy when clients share their local parity datasets with the server
 - Analyzed the convergence rate and iteration complexity of CodedFedL, by treating CodedFedL as a stochastic gradient descent algorithm
 - In numerical experiments using practical network parameters and benchmark datasets, CodedFedL provided gains of up to $15\times$ in comparison to benchmark schemes

International Visiting Student

Summer 2015

IUSSTF-Viterbi Program, Los Angeles, CA.

Mentor – [Prof. Salman Avestimehr](#)

- Towards Faster Algorithms for Processing Large Data on Graphs
 - Studied spectral graph theory and its application in signal processing of graph data – cut-off frequency, optimal sampling and bandlimited interpolation
 - Explored existing semi-supervised and active learning methods for data on graphs
 - Proposed a Random Jump model based on Graph Laplacian for sampling with low time complexity

Undergraduate Research Assistant

2013-2016

Multimedia Wireless Networks Group, IIT Kanpur, Kanpur, India

Mentor – [Prof. Aditya K. Jagannatham](#)

- Scheduling for Efficient Utilization of Time Resource in Wireless Networks
 - Worked on the problem of user scheduling for efficient wireless resource utilization, under resource allocation fairness constraints
 - Proposed two opportunistic schemes for scheduling users in a time slotted system with wireless Rayleigh-fading channel
 - Simulations predicted stochastically improved performance compared to Round Robin scheme alongside satisfaction of any arbitrary time resource allocation fairness constraints

Undergraduate Research Intern

Summer 2013

Summer Undergraduate Research Grant for Excellence (SURGE), IIT Kanpur, Kanpur, India.

Mentor – [Prof. Aditya K. Jagannatham](#)

	<ul style="list-style-type: none"> • Channel Estimation and Capacity in MIMO Wireless Communication Systems <ul style="list-style-type: none"> – Analyzed capacity lower bound for a MIMO system obeying Block-Fading law using LMMSE estimator for channel estimation at the receiver – Obtained an optimal number of transmit antennas for optimizing the capacity lower bound 	
SELECTED TALKS	<ul style="list-style-type: none"> • TEE-GPU Cooperative Learning: Privacy and Security Without the Price (with Hanieh Hashemi) <i>Presentation, Enclaved AI/ML Workshop, Private AI Research Institute</i> • Federated deep learning: On-device learning for CV and NLP (with Chaoyang He) <i>Finalist Team Presentation, Qualcomm Innovation Fellowship</i> • Trustworthy and Scalable Federated Learning (with Salman Avestimehr, Roushdy Elkordy, Chaoyang He, Jinhyun So) <i>CCF Advanced Disciplines Lecture, Institute of Computing Technology, Chinese Academy of Sciences</i> • Coded Computing for Federated Learning at the Edge <i>Presentation, FL Workshop on User Privacy and Data Confidentiality, ICML</i> 	<p><i>Jul 2021</i></p> <p><i>May 2021</i></p> <p><i>Apr 2021</i></p> <p><i>Jul 2020</i></p>
COMMUNITY SERVICE	<p>Invited Journal Reviewer</p> <ul style="list-style-type: none"> • IEEE Journal on Selected Areas in Communications • IEEE Transactions on Information Theory • IEEE Journal on Selected Areas in Information Theory • IEEE Transactions on Communications <p>Invited Conference/Workshop Reviewer</p> <ul style="list-style-type: none"> • IEEE International Symposium on Information Theory • IEEE Information Theory Workshop <p>Mentorship</p> <ul style="list-style-type: none"> • Mentor, Graduate Application Mentorship Program (GradAMP), USC • Student Guide, Counselling Service, IIT Kanpur • Secretary, Fine Arts Club, IIT Kanpur <p>Others</p> <ul style="list-style-type: none"> • Ambassador Caller, Alumni Contact Program, IIT Kanpur • Secretary, Fine Arts Club, IIT Kanpur 	<p><i>2017-2021</i></p> <p><i>2017-2021</i></p> <p><i>2021</i></p> <p><i>2013-2016</i></p> <p><i>2013-2014</i></p> <p><i>2013-2015</i></p> <p><i>2013-2014</i></p>
SELECTED COURSEWORK	<p>Algorithms and Artificial Intelligence</p> <ul style="list-style-type: none"> • Deep Learning • Data Structure and Algorithms • Fundamentals of Computing • Machine Learning for Computer Vision • Artificial Intelligence Programming <p>Mathematics</p> <ul style="list-style-type: none"> • Real Analysis • Probability and Statistics • Linear Algebra • Complex Analysis • Partial Differential Equations • Random Processes in Engineering <p>Communication and Signal Processing</p> <ul style="list-style-type: none"> • Error Correcting Codes • Wireless Communications • Information Theory • Topics in Cryptography and Coding • Digital Signal Processing • Digital Communication Networks 	

TECHNICAL
SKILLS

Python (including PyTorch and Keras APIs), Amazon Web Services, C, MATLAB